



DeepLearning.AI

# Module 2 introduction

---

## Information Retrieval Foundations



## Prompts

Unstructured  
Conversational



## Retriever

Needs to rapidly find  
relevant documents  
despite this mess!



## Documents

Wide range of formats  
Designed for humans to read

# Module Topics

- Widely used retrieval techniques
- Theoretical foundations and relative strengths of each technique, plus how they're used in combination
- Evaluation strategies
- Hands-on examples and a programming assignment

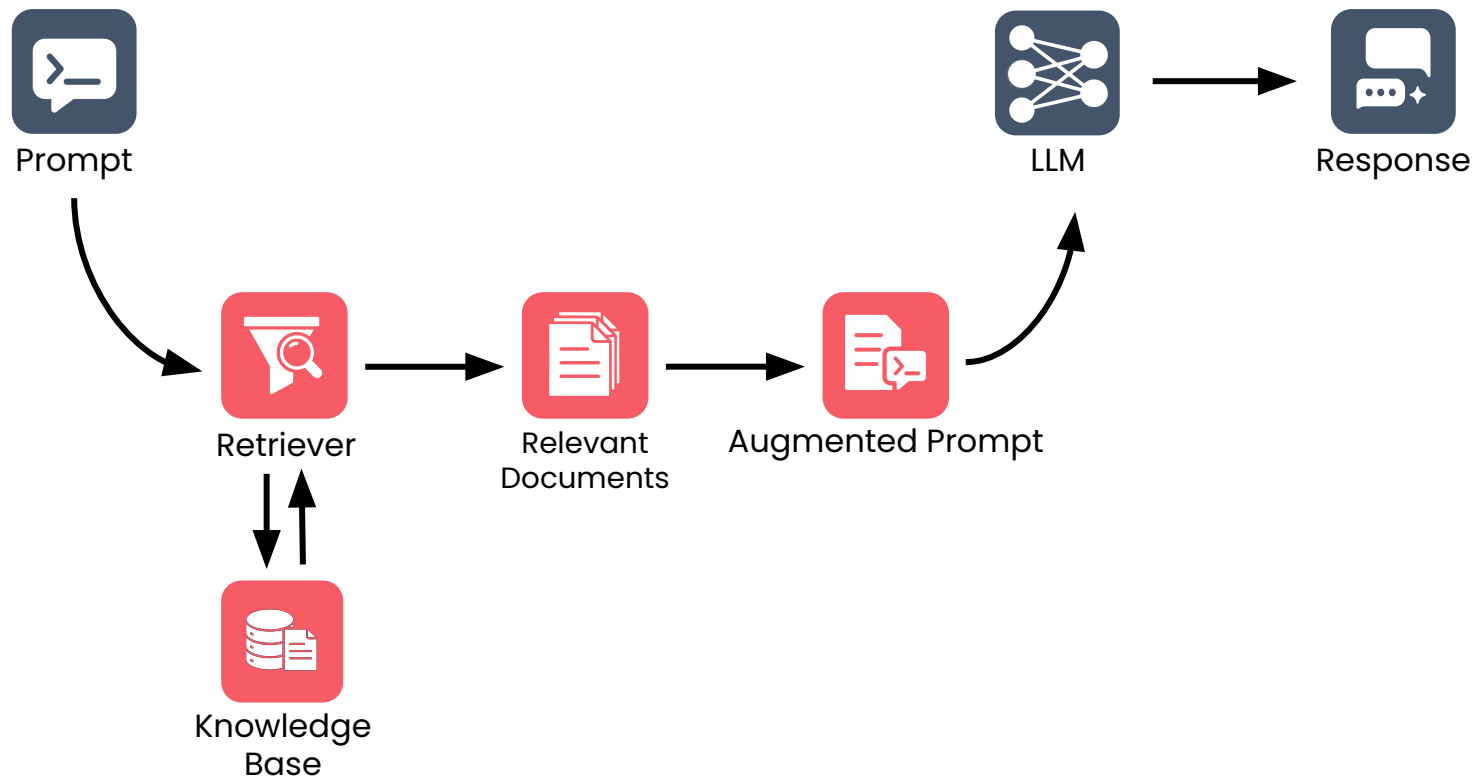


DeepLearning.AI

# Retriever architecture overview

---

## Information Retrieval Foundations




# Two Search Approaches



## Keyword Search

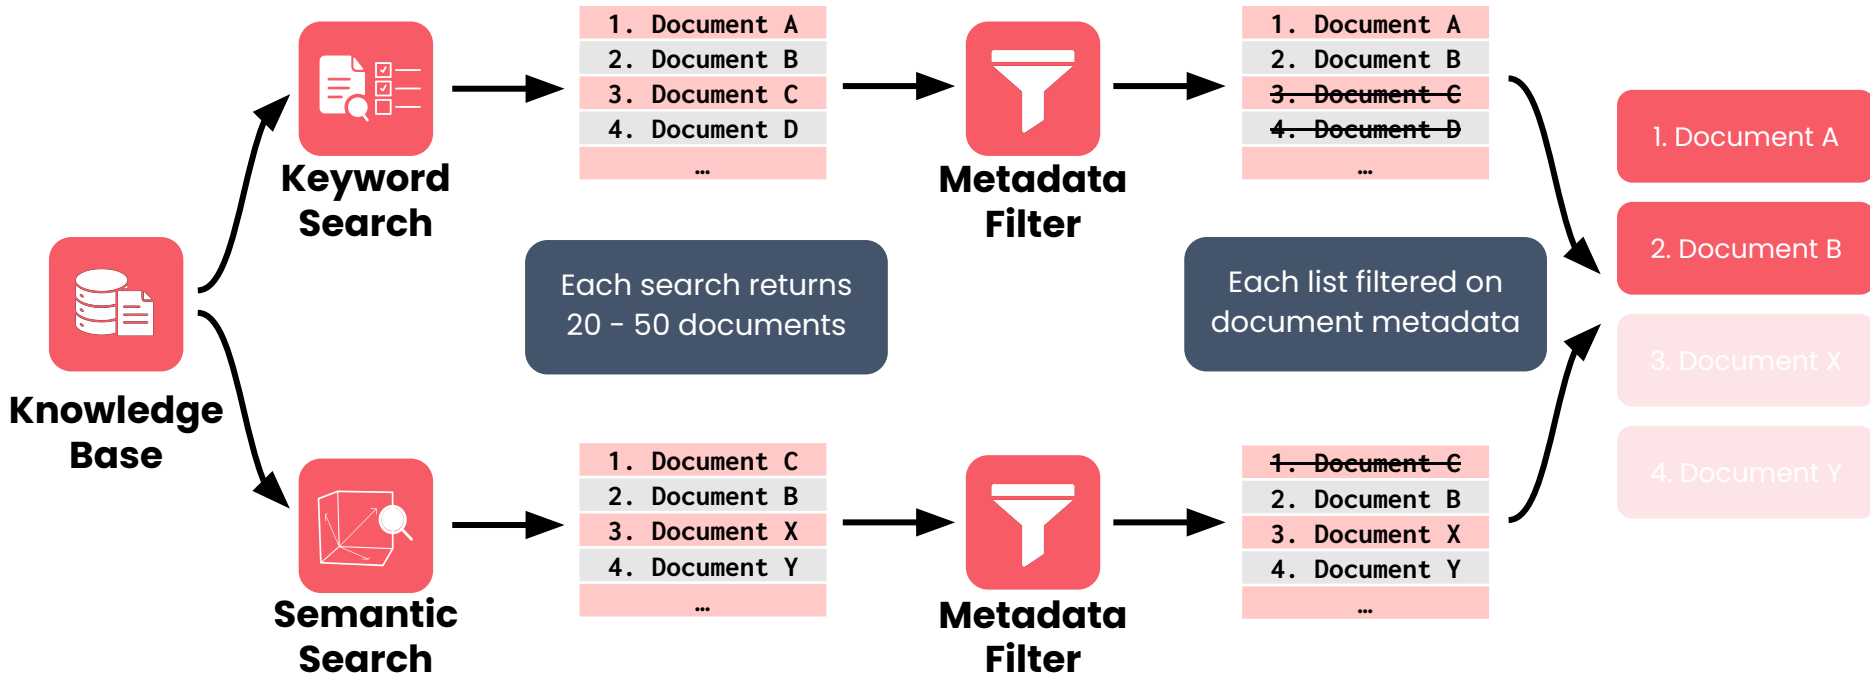
Looks for documents containing the **exact words** found in the prompt.



## Semantic Search

Looks for documents with **similar meaning** to the prompt.

# Search Techniques



# Hybrid Search



## **Keyword Search**

Ensures sensitivity to exact words the user included in the prompt



## **Semantic Search**

Finds documents with similar meaning, even without matching words



## **Metadata Filtering**

Excludes documents based on rigid criteria



High-performing retrievers balance all three techniques based on project needs.





DeepLearning.AI

# Metadata filtering

---

Information Retrieval  
Foundations

# Metadata Filtering

Uses rigid criteria to narrow down documents based on metadata like title, author, creation date, access privileges, and more.



## The Daily Maple



**Title**



**Publication  
Date**



**Author**



**Section**




**Tags**

```
SELECT * FROM articles  
WHERE publication_date = '2023-10-01';
```

Return all articles published on specific date

```
section = "Opinion"  
author = "Michael Chen"  
date = June to July 2024
```

# Metadata Filtering In RAG



Metadata filtering doesn't perform retrieval, it **narrows down results from other techniques** based on user attributes, not query content.

## Subscription Filtering



Filter: Exclude all articles with metadata "subscription = paid"

## Geographic Filtering



Filter: Include only articles with metadata "region = North America"

# Advantages and Limitations

## Pros

- Simple to understand and debug
- Fast, optimized, mature, and reliable
- Enforces strict retrieval rules, matching exact filter criteria

## Cons

- Not true search
- Rigid, ignores content, and provides no way for ranking
- Useless alone




DeepLearning.AI

# Keyword search - TF-IDF

---

## Keyword Search TF-IDF

# Introduction to Keyword Search



## Keyword Search



### **Decades of proven simplicity**

The simplicity and effectiveness that has powered retrieval for decades makes keyword search a key component of modern RAG systems.



How can I make New York style pizza at home?



**Keyword  
Search**



Doc 1

Use bread flour for New York pizza dough.



Doc 2

New York pizzerias stay open late for home delivery.



Doc 3

At New York Pizza Mexico, they serve pizza with jalapeño

# Bag of Words

Word order is ignored, only word presence and frequency matter

Prompt

**"Making pizza without a pizza oven"**

Keywords

pizza

2

without

1

making

1

oven

1

a

1

**Bag of words**



# Sparse Vectors

“Making pizza without a pizza oven”

making  
1

pizza  
1

cake  
0

pan  
0

without  
1

drink  
0

pasta  
0

a  
1

oven  
1

pan  
0

pasta  
0

burger  
0

taco  
0

with  
0

salad  
0

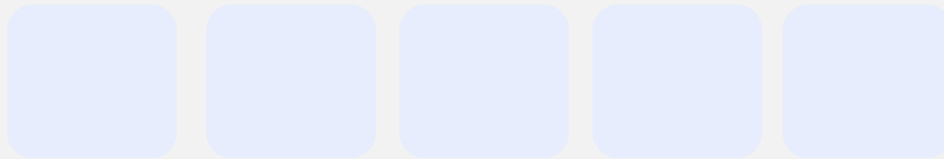
blend  
0

tea  
0

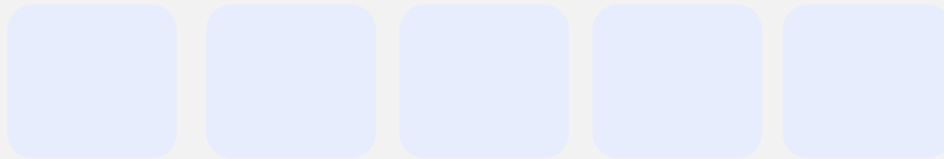
the  
0

Most words aren't used. The bag of words is sparse, with few non-zero entries.

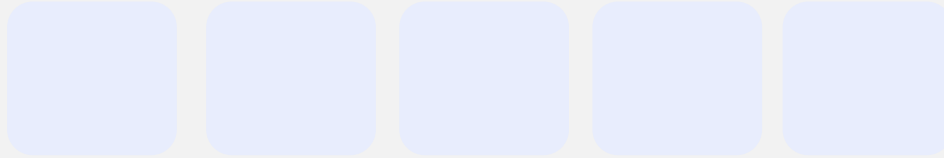
**Document 1**



















**Document 2**



**Document 3**



	Doc 1	Doc 2	Doc 3	Doc 4	Doc 5
<b>making</b>					
<b>without</b>					
<b>oven</b>					



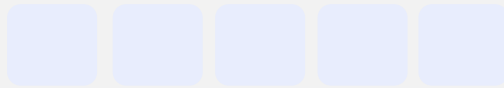


**Making pizza without a  
pizza oven**



**Retriever**

**Prompt  
sparse  
vector**



	Doc 1	Doc 2	Doc 3	Doc 4	Doc 5
making					
without					
oven					

# "Making pizza without a pizza oven"

	Doc 1	Doc 2	Doc 3	Doc 4	Doc 5
<b>making</b>					
<b>pizza</b>					
<b>without</b>					
<b>a</b>					
<b>oven</b>					

## "Making pizza without a pizza oven"

	Doc 1	Doc 2	Doc 3	Doc 4	Doc 5
<b>making</b>	+1			+1	+1
<b>pizza</b>					
<b>without</b>					
<b>a</b>					
<b>oven</b>					

# "Making pizza without a pizza oven"

	5/5 Doc 1	3/5 Doc 2	1/5 Doc 3	4/5 Doc 4	4/5 Doc 5
making	+1			+1	+1
pizza	+1		+1		
without	+1	+1		+1	+1
a	+1	+1		+1	+1
oven	+1	+1		+1	+1



# Frequency Based Scoring

Example Query: pizza oven

## Document 1

Homemade **pizza** in **oven** is better  
than frozen **pizza**

Contains: Pizza (2x) Oven (1x)

**Simple Scoring = 2 points**

**TF Scoring = 3 points**

## Document 2

Wood-fired **oven** is a better **oven** than  
a stone **oven** for cooking **pizza**

Contains: Pizza (1x) Oven (3x)

**Simple Scoring = 2 points**

**TF Scoring = 4 point**

# Normalized TF Scoring

Longer documents may contain **keywords many times** simply because they are **longer**.

**Solution: Normalize by document length**

**Score** = (Number of keyword occurrences) / (Total words in document)

# TF-IDF

Basic TF scoring treats **all words equally**, whether they're common filler words or rare, meaningful terms.

**Solution: Weight terms using “inverse document frequency” (IDF).**

$$\textbf{Score} = \text{TF}(\text{word}, \text{doc}) \times \log(\text{Total docs} / \text{Docs containing word})$$

1

## Count Documents

For each word

making

pizza



docs word appears in

---

total docs

1

## Count Documents



**Pizza**

Appears in 5 out of 100 documents

**DF** = 5/100

0.05



**The**

Appears in all 100 documents

**DF** = 100/100

1.0

1

## Count Documents



**Pizza**

Appears in 5 out of 100 documents

$$DF = 5/100$$

0.05



**The**

Appears in all 100 documents

$$DF = 100/100$$

1.0

2

## Flip to reward rare words



**Pizza**

High Score

$$IDF = 1/0.05$$

20



**The**

Low Score

$$DF = 1/1.0$$

1.0

1

## Count Documents



Pizza

Appears in 5 out of 100 documents

$$DF = 5/100$$

0.05



The

Appears in all 100 documents

$$DF = 100/100$$

1.0

2

## Flip to reward rare words



Pizza

High Score

$$IDF = 1/0.05$$

20



The

Low Score

$$DF = 1/1.0$$

1.0

3

## Apply log



Pizza

Still higher

$$\log(1/20)$$

1.30



The

Too common, no weight

$$\log(1)$$

0

## TF

- **making:** 0.4 (appears in 3/5 docs)  
moderately common
- **pizza:** 0.7 (appears in 2/5 docs)  
less common
- **without:** 0.2 (appears in 4/5 docs)  
very common
- **a:** 0.1 (appears in 4/5 docs)  
very common
- **oven:** 0.2 (appears in 4/5 docs)  
very common

## TF-IDF

	Doc 1	Doc 2	Doc 3	Doc 4	Doc 5
making	0.4			0.4	0.4
pizza	0.7		0.7		
without	0.2	0.2		0.2	0.2
a	0.1	0.1		0.1	0.1
oven	0.2	0.2		0.2	0.2



# “Making pizza without a pizza oven”



## High Scoring Document

Pizza  
oven

1.3

1.1



## Low Scoring Document

a  
without

0

0.1

Documents with **rare keywords** score **higher** than documents with common words



Modern systems use a slightly refined version called  
BM25



DeepLearning.AI

# Keyword search – BM25

---

Information Retrieval  
Foundations

# BM25 Scoring

BM25 (Best Matching 25) was named as the 25th variant in a series of scoring functions proposed by its creators.

$$IDF * \frac{TF * (k_1 + 1)}{TF + k_1 * (1 - b + b * (\frac{\text{document length}}{\text{average document length}}))}$$

- This gives the score for a **single keyword**
- **Sum scores** across all keywords for total relevance score for a document

## 1 Term Frequency Saturation



TF-IDF

"pizza" 10 times = Score X  
"pizza" 20 times = Score 2X



BM25

"pizza" 10 times = Score X  
"pizza" 20 times = Score 1.3X

Term Frequency Saturation

## 2 Document Length Normalization



TF-IDF

Short Doc = Good Score  
Long Doc = Heavy Penalty

Too aggressive



BM25

Short Doc = Good Score  
Long Doc = Smaller Penalty

Document length  
normalization

# BM25 Tunable Parameters


## $k_1$ – Term Frequency Saturation

- **Controls:** How much term frequency influences the score.
- **Range:** Typically between 1.2 and 2.0.
- **Effect:** Higher values increase the impact of term frequency; lower values reduce it.

## $b$ – Length Normalization

- **Controls:** The degree of normalization for document length.
- **Range:** Between 0 (no normalization) and 1 (full normalization).
- **Effect:** Balances favoring shorter vs. longer documents.

# TF-IDF vs BM25



**BM25** = Standard keyword search algorithm in production retrievers

**better performance + same cost + more flexibility**  
than TF-IDF

# Keyword Search Overview



Match documents by **keyword frequency**



## TF-IDF

- Keyword rarity
- Term frequency
- Document length

## BM25

Most commonly used


- Document length normalization
- Term Frequency Saturation



# Keyword Search Strengths



**Simplicity**



**Guaranteed keyword  
matching**



DeepLearning.AI

# Semantic search introduction

---

Information Retrieval  
Foundations


In order to improve on lexical search it's necessary to capture not only the presence of words, but their **meaning**.

### Cannot match synonyms

"happy"

"glad"

### Incorrectly matches different meanings



# Semantic Search vs. Keyword Search

- Prompt and documents each get a vector
- Vectors compared to generate scores
- The main difference is how vectors are assigned
  - **Keyword Search:** count words
  - **Semantic Search:** use embedding model


# Understanding Embedding Models

Embedding models map tokens, to a location in space. This location is represented by a **vector**.

**"Pizza"** → vector  $[3, 1]$

**"Bear"** → vector  $[5, 2]$

In two dimensions, these can be represented as points




**Vector Space**

# Understanding Embedding Models

No simple interpretation of X  
and Y axis


...instead...

points “float around” in  
space and similar words  
cluster together




**Vector Space**


# Understanding Embedding Models



2 dimensions



3 dimensions



100 - 1,000+ dimensions

More dimensions means more room  
to form clusters and capture  
nuanced relationships


Same principles hold  
Close vectors, similar meanings

## Individual Words

cat

happy

Word  
Embedding  
Model




## Sentences

The weather is nice

What a lovely day


Sentence  
Embedding  
Model



## Documents



Document  
Embedding  
Model






# Sentence Embedding Example


1 He spoke softly in class

2 He whispered quietly during class

3 Her daughter brightened the gloomy day




# Measuring Vector Distance



**Euclidean Distance**

Measures how far apart two vectors are by drawing a straight line from one vector to the other – the shortest possible distance between them.

# Measuring Vector Distance



**Cosine Similarity**

Looks at the similarity in the **direction** of two vectors, regardless of whether they're close to one another in space

-1

Opposite  
direction


0

Perpendicular

1

Same  
direction

# Measuring Vector Distance



**Dot product**

Measures the length of the projection of one vector onto another.

**Negative**

Opposite  
direction


**0**

Perpendicular

**Positive**

Same  
direction

# Measuring Vector Distance



**Cosine Similarity**

-1

Opposite  
direction


0

Perpendicular

1

Same  
direction

Higher values, closer vectors



**Dot product**

Negative

Opposite  
direction


0

Perpendicular

Positive

Same  
direction

# Semantic Search






DeepLearning.AI

# Semantic search embedding model deepdive

---

## Information Retrieval Foundations


# Embedding Models




How can an embedding model know to place similar text together, and dissimilar text farther apart?



# Positive and Negative Examples in Training




**Positive Pair**



**Negative Pair**


# Positive and Negative Examples in Training




**Compile massive training dataset  
of positive and negative pairs**

# Initial Random Vectors in Embedding Models

## Random Initialization



# Contrastive Training Process



Look at where each pair was placed in vector space




Scores better when **closer together**



Scores better when **farther apart**


**Score Results**

# Contrastive Training Process




# Contrastive Training Process


- Update internal parameters based on scoring the positive and negative pairs
- Repeat the process: Embed  $\rightarrow$  Score with Pairs  $\rightarrow$  Update parameters
- Iteratively repeat the process, improving the model



Early in training



Mid-training




Trained model

# Contrastive Training Process


**Beginning of Training**  
Random Positions



**Training**  
Pushing and pulling



**After Training**  
Meaningful Embeddings




● "He could smell the roses"

● "A field of fragrant flowers"

● "The lion roared majestically"

# Scaling up Contrastive Learning

- In reality every vector is simultaneously pushed and pulled in many directions
- Using 100s or 1,000s of dimensions creates more space in which to push and pull vectors
- Eventually vectors pulled near similar words or text





# Key Takeaways

- Semantic vectors are abstract and somewhat random
- Before training: locations in space have no meaning
- After training: locations have meaning because clusters of similar text have formed
- Only compare vectors from same embedding model



DeepLearning.AI

# Vector embeddings in RAG

---

Hybrid Search in  
Information Retrieval

# Key Strategies



## **Metadata Filtering**

Uses rigid criteria stored in document metadata to narrow down search results  
Fast, easy, yes-no filter, but can't be used alone



## **Keyword Search**

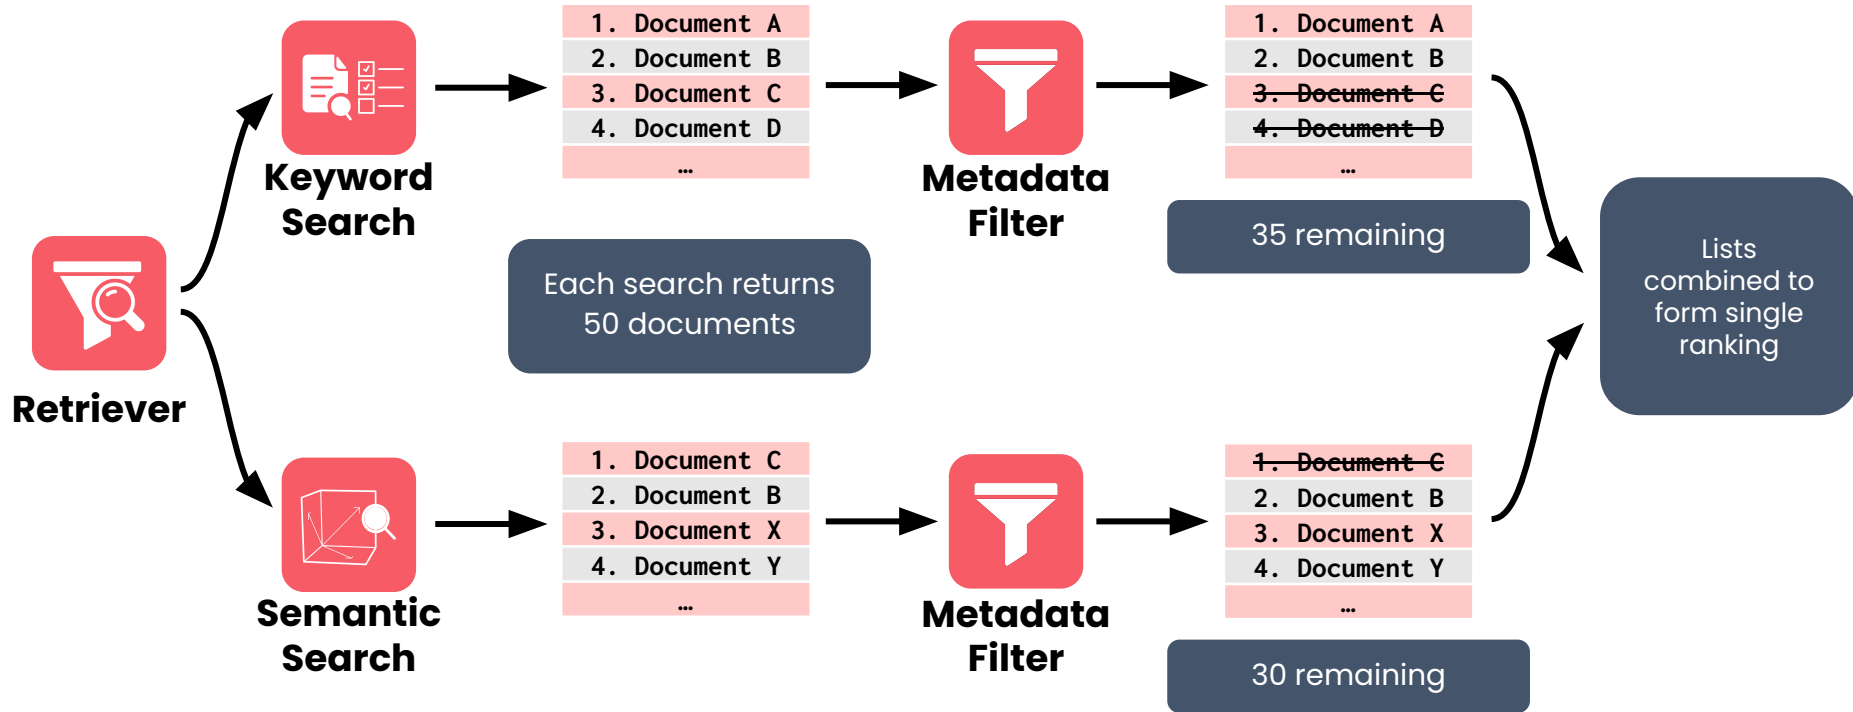
Scores documents based on having the same keywords found in the prompt  
Fast, performs especially well when keywords matter, but relies on exact matches



## **Semantic Search**

Scores and ranks documents based on having similar meaning to the prompt  
Slower, computationally expensive, but more flexible

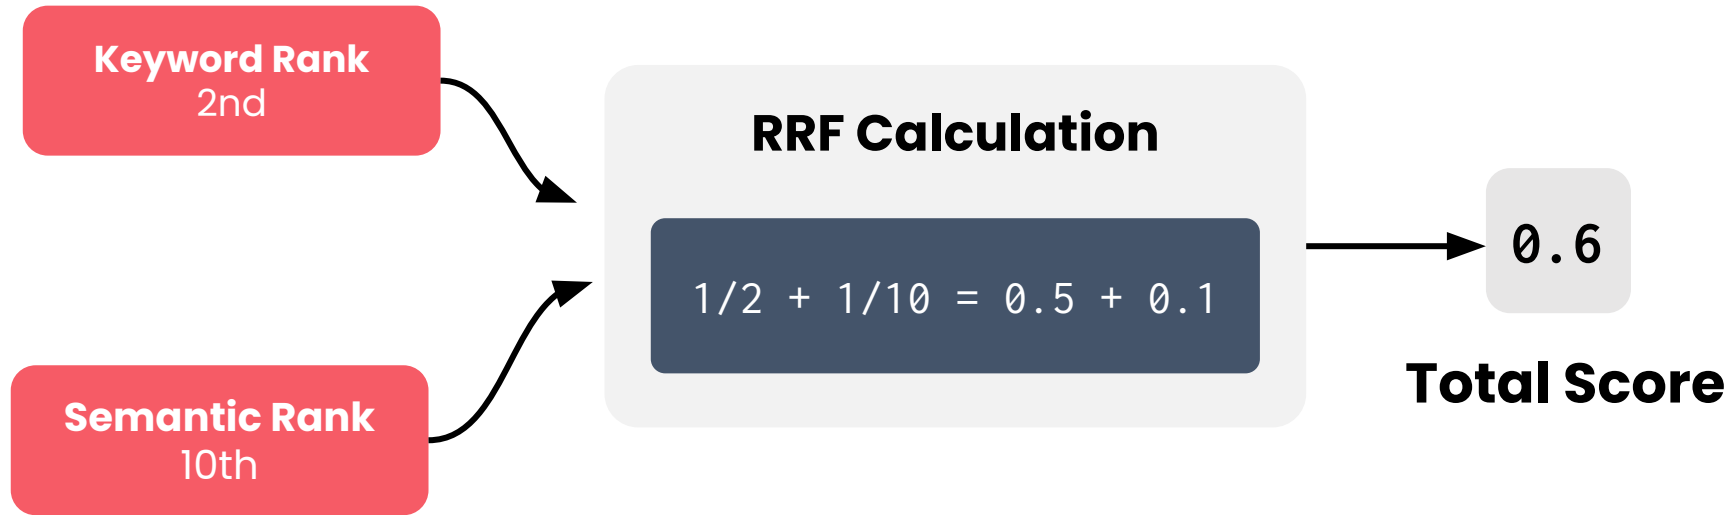
# Hybrid Search



# Reciprocal Rank Fusion

- Rewards documents for being highly ranked on each list
- Control weight of keyword vs. semantic ranking
- Score points equal to reciprocal of ranking  
1st = 1 point, 2nd = 0.5 points, etc.
- Total points from all ranked list used to perform final ranking

$$\frac{1}{k + \text{rank in list 1}} + \frac{1}{k + \text{rank in list 2}} + \dots + \frac{1}{k + \text{rank in list n}}$$



# Reciprocal Rank Fusion

$$\frac{1}{k + \text{rank in list 1}} + \frac{1}{k + \text{rank in list 2}} + \dots + \frac{1}{k + \text{rank in list n}}$$

**When  $k = 0$**

Top ranked document shoots to top of overall ranking  
**1st vs 10th: 10x difference**


**When  $k = 50$**

Single high rank doesn't dominate overall ranking  
**1st vs 10th: 1.2x difference**


**RRF only cares about ranks, not scores**

# Beta: Weighting Semantic vs. Keyword

$\beta = 0.8$



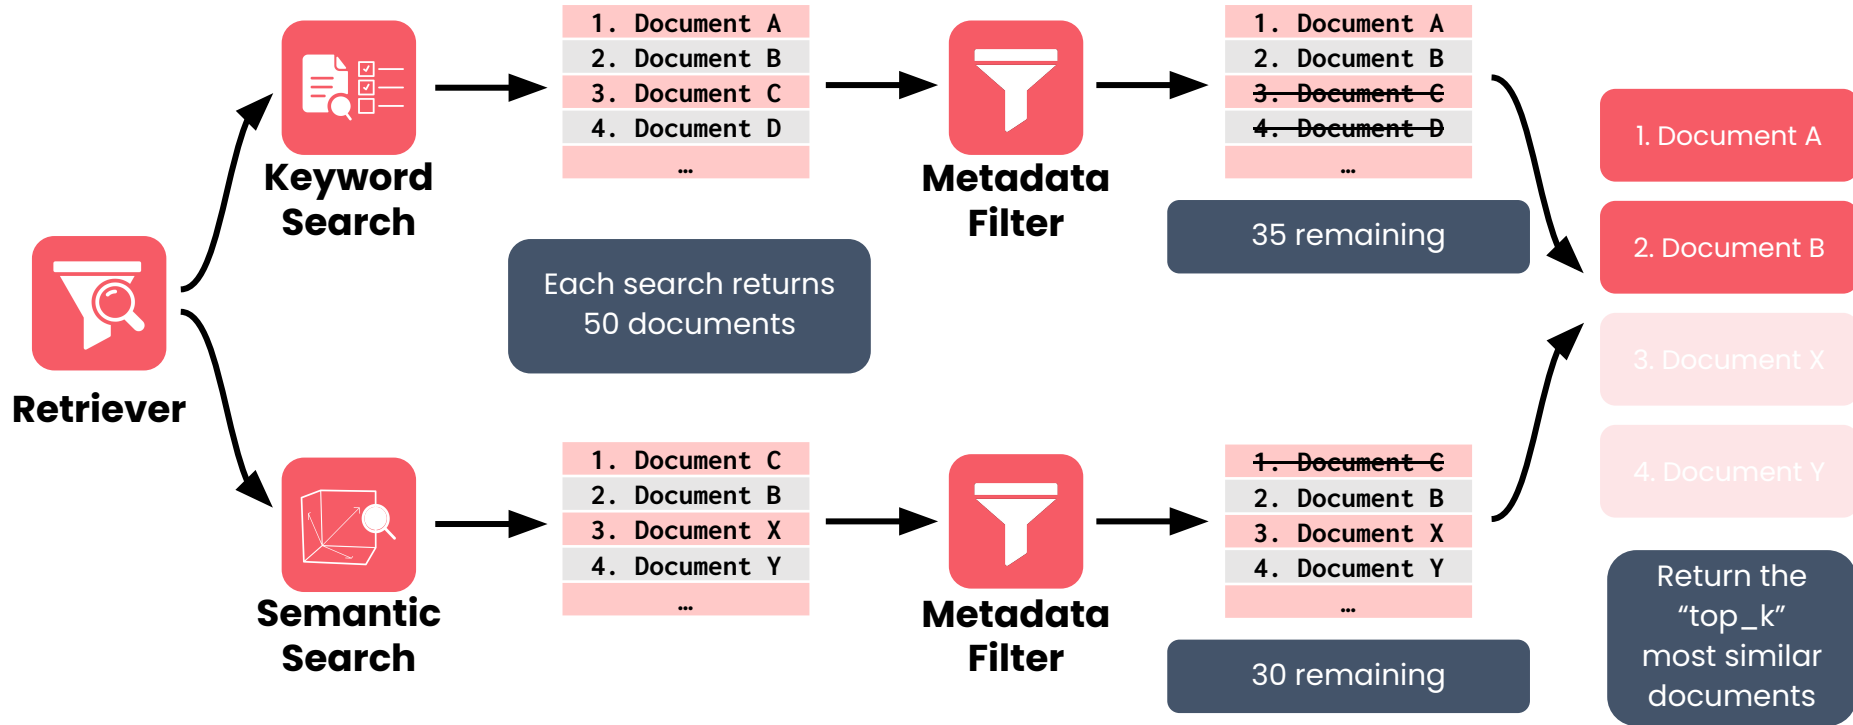
$\beta = 0.7$



**If exact keyword matching is important, set a lower beta**



# Hybrid Search





DeepLearning.AI

# Hybrid search

---

Information Retrieval  
Foundations

# Retrieval Quality Metrics

Common ingredients to most retriever quality metrics:

## **The Prompt**

The specific prompt being evaluated

## **Ranked Results**

Documents returned in ranked order

## **Ground Truth**

All documents labeled as relevant or irrelevant

If you want to evaluate your retriever  
**you need to know the correct answers**

# Precision and Recall

## **Precision**

Measures how many of the returned documents are relevant

$$\text{Relevant Retrieved} / \text{Total Retrieved}$$

---

## **Recall**

Measures how many of the relevant documents are returned

$$\text{Relevant Retrieved} / \text{Total Relevant}$$

# Example



**10 Relevant Documents** in Knowledge Base



## First Run

Retrieved: 12 Documents

Relevant: 8 Documents

**Precision** (8/12)

66%

**Recall** (8/10)

80%



## Second Run

Retrieved: 15 Documents

Relevant: 9 Documents

**Precision** (9/15)

60%



**Recall** (9/10)

90%




**Precision** penalizes for **returning irrelevant** documents

**Recall** penalizes for **leaving out relevant** documents

**100% Precision & Recall:** Rank relevant documents most highly and only return those

# Top k



Top K

- Retrieval metrics are influenced by **how many documents the retriever returns**
- Metrics are discussed in terms of **top-k** documents

# Example

Rank	Relevant
1	relevant
2	-
3	-
4	relevant
5	-
6	relevant
7	relevant
8	-
9	relevant
10	relevant

**Precision @5**

2 out of 5

40%

**Precision @10**

6 out of 10

60%

**Recall@10**

6 out of 8 total  
relevant

75%

**Top-5 or Top-1 is stricter**

**Top-5 to Top-15 often used**

# Mean Average Precision

**MAP@K** evaluates average precision for relevant documents in first K documents. It is built off a related metric called “average precision”.

Rank	Item	Precision@K
1	relevant	1/1 1.0
2	-	1/2 0.5
3	-	1/3 0.3
4	relevant	2/4 0.5
5	relevant	3/5 0.6
6	-	3/6 0.5

Rewards ranking relevant documents highly

1

**Sum precisions** for relevant docs only.

$$1 + 0.5 + 0.6 = 2.1$$

2

**Divide** by number of relevant documents

$$2.1 / 3 = 0.7$$

3

This calculation gives **Average Precision** or AP, for Mean Average Precision you find the average AP value across many prompts



# Reciprocal rank

Measures the rank of the first relevant document in the returned list

$$\text{Reciprocal Rank} = 1 / \text{Rank}$$

**First relevant at rank 1**

1.0

**First relevant at rank 2**

0.5

**First relevant at rank 4**

0.25

The later the first relevant document appears, the worse the reciprocal rank

Mean Reciprocal Rank (MRR) averages over many prompt

# Mean Reciprocal Rank

Search 1 **First relevant at rank 1**

1.0

Search 2 **First relevant at rank 3**

0.33

Search 3 **First relevant at rank 6**

0.17

Search 4 **First relevant at rank 2**

0.5

1 **Sum all ranks**

$$1.0 + 0.33 + 0.17 + 0.5 = 2.0$$

2 **Divide by number of searches**

$$2.0 / 4$$

**MRR = 0.5**

# How to use retriever metrics

## Recall or recall@K

Most cited metric, captures fundamental goal of finding relevant documents

## Precision & MAP

Asses irrelevant documents and ranking effectiveness

## Mean Reciprocal Rank

How well model performs at the very top of ranking

## Metrics help:

- Evaluate retriever performance
- Check if adjustments improve results

**All metrics depend on having  
ground truth relevant  
documents**



DeepLearning.AI

# Module 2 Conclusion

---

Information Retrieval  
Foundations

# Conclusion

- **Keyword Search**  
Ranks by keyword frequency exact matches
- **Semantic Search**  
Ranks by meaning, flexible
- **Metadata Filtering**  
Excludes by criteria
- **Hybrid Search**  
Combines all three techniques

## Evaluation Metrics

Precision & Recall

MAP

Mean Reciprocal Rank

Measure improvement from adjusting tunable parameters in hybrid search